

Eyal Orbach - id 015369317

Daniel Juravski - id 206082323

- **how large were the training and test sets?**

The training and test sets were both of size 1000 (500 positive 500 negative)
The training set was divided to 80% train 20 percent validation.

- **did your network succeed in distinguishing the two languages?**

yes

- **how long did it take (both wall-clock time (i.e., number of seconds), and number of iterations)?**

It runs for 5 epoch, 16 seconds.

validation was 100% from epoch 2

each epoch runs over the 800 training samples and 200 validation samples

- **did it succeed only on the train and not on the test?**

succeeds on both

- **what you did in order to make it work?**

These are the parameters we used

LAYERS = 1 # layers of the lstm

INPUT_DIM = 16 # size of the embedding vector

HIDDEN_DIM = 20 # size of the state vector

LSTM_OUTPUT_SIZE = 10 # size of y, the input to the MLP

N1 = 8 # size of hidden layer inside the MLP

we used RMSPropTrainer as the optimizer, which converged fastest.

